

# MAF: Masked Autoregressive Flow for Density Estimation

Lorne Whiteway  
lorne.whiteway@star.ucl.ac.uk

Astrophysics Group  
Department of Physics and Astronomy  
University College London

27 April 2020

Find the presentation at <https://tinyurl.com/???>

# Context

- ▶ *Supervised machine learning* i.e. selecting from a highly-parameterised family of non-linear functions to fit some training data.
- ▶ Fitting is done to optimise a utility function that **rewards a close match to the training data** and **penalises complexity** i.e. *enforces regularization*.
- ▶ Regularisation is also provided by any lack-of-flexibility in the fitting functions.
- ▶ Bayesian framework: **adherence to training data = likelihood**; **regularisation = prior**.

# Density Estimation

- ▶ We want to do ML to do density estimation:
- ▶ Find the probability density  $p$ , given a set of training data  $\{x_i\}$  that we suppose to be a fair random sample from  $p$ .

- ▶ We will discuss the MAF algorithm: Papamakarios et al. 2018; Masked Autoregressive Flow for Density Estimation; <https://arxiv.org/abs/1705.07057>.
- ▶ Based on earlier algorithms of which the most relevant is MADE: Germain et al. 2015; MADE: Masked Autoencoder for Distribution Estimation; <https://arxiv.org/abs/1502.03509>.
- ▶ I will focus first on MADE as all the key ideas are present there.

# Analogy

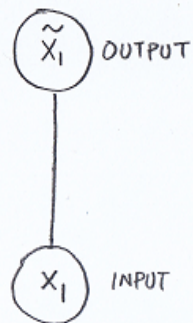
- ▶ An *autoencoder* is like a parrot - it learns to mimic.
- ▶ An *autoregressive autoencoder* is like a deaf parrot - it learns to play the probabilities (and hence becomes a density estimator).

# Parrot training

- ▶ Simple example: single binary outcome. We want the parrot to mimic us when we say 'Yes' or 'No'.
- ▶ So give lots of 'Yes', 'No' training data; give the parrot treats but reduce the treats by the non-positive number  $\log(\% \text{ correctness of the parrot's response})$ . This is called *cross-entropy loss*.

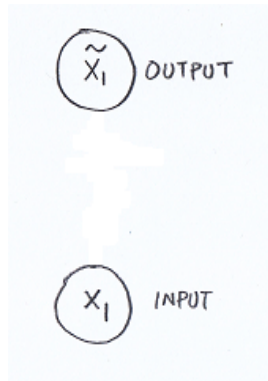
# Parrot training

- ▶ The neural network is incentivised to model the identity function  $\tilde{x}_1 = x_1$ .
- ▶ Depending on the flexibility of the neural network the parrot will get this exactly or approximately correct.



# Deaf parrot

- ▶ Deaf parrot = no link between input and output.
- ▶ Parrot still squawks and gets feedback (more treats if the squawk accidentally matched the unheard training input).
- ▶ If say 75% of the training data was 'Yes' then the parrot will learn to say 'Yes' more often. Optimal strategy will be a word that is 75% 'Yes' and 25% 'No'.





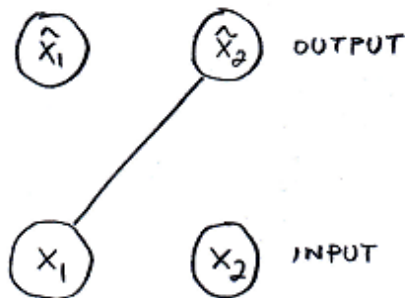
- ▶ So by removing links from the neural network but still rewarding the network for mimicry we create a density estimator.

## Two word phrase

- ▶ Now consider two word phrases ("Yes, No").
- ▶ If there is no correlation between the two words then we just need two deaf parrots.
- ▶ But generally  $p(x_1 \& x_2) = p(x_1)p(x_2|x_1)$ .
- ▶ To model this: one deaf parrot for  $p(x_1)$  as before. But for the second parrot: unmask its ears, let it hear the training word told to the first parrot, the re-mask its ears for the second word.
- ▶ Give treats as before (still encouraging mimicry).
- ▶ Easy to show that the second parrot learns  $p(x_2|x_1)$ .

## Two word phrase

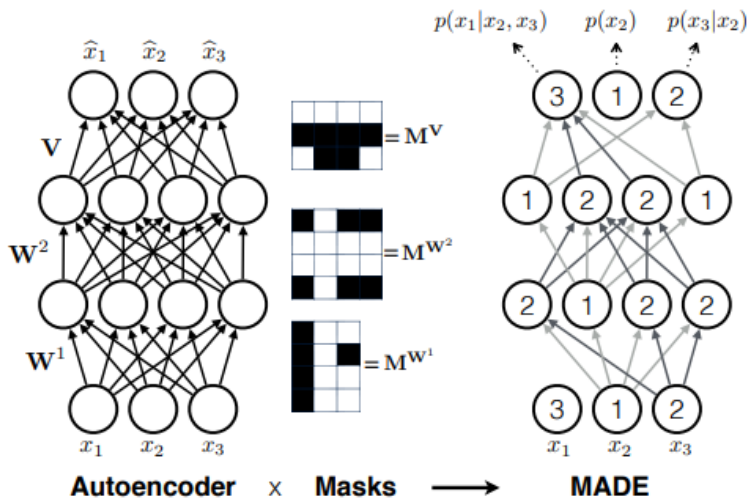
- More generally if  $\tilde{x}_i$  is linked only to  $x_1, \dots, x_{i-1}$  then we learn  $p(x_1), \dots, p(x_i | x_1 \dots x_{i-1})$  whose product is  $p(x)$ .



# Autoencoder; Autoregressive autoencoder

- ▶ The 'hearing' parrot (that mimics) is an *autoencoder*.
- ▶ The network of partially-deaf parrots (that learns the conditional probabilities) is called an *autoregressive autoencoder*.
- ▶ Autoregression appears in other contexts e.g. predicting the next element in a causal sequence.

Figure 1 from MADE paper



# Sampling from the inferred probability density

- ▶ Use  $(0, 0)$  as input to get  $(\tilde{x}_1, \tilde{x}_2)$  as output;
- ▶ Randomly choose  $s_1$  in  $\{0, 1\}$  with probability  $\tilde{x}_1$  of being 1;
- ▶ Use  $(s_1, 0)$  as input to get new  $(\tilde{\tilde{x}}_1, \tilde{\tilde{x}}_2)$  as output;
- ▶ Randomly choose  $s_2$  in  $\{0, 1\}$  with probability  $\tilde{\tilde{x}}_2$  of being 1;
- ▶ Then  $(s_1, s_2)$  is a new sample from the inferred density.
- ▶ Takes  $D$  evaluations of the net to get one new sample.

# More complicated distributions

- ▶ So far we have looked at binary distributions.
- ▶ Here the probability density could be described by a single number ( $= p(x_i = 1)$ ). So output  $\tilde{x}_i$  has dual meaning as 'best guess at  $x_i$ ' and as ' $p(x_i = 1|x_1, \dots, x_{i-1})$ '.

# More complicated distributions

- ▶ To handle continuous distributions:
- ▶ Assume each univariate conditional distribution  $p(x_i|x_1, \dots, x_{i-1})$  is Gaussian;
- ▶ Two output cells for each input: one for mean and one for (log of) standard deviation - from this we easily calculate the conditional probabilities and hence the overall probability of any given input;
- ▶ Loss function is now sum of negative logs of probabilities of training data.





# Sampling from the inferred probability density $\tilde{p}$

- ▶ As before we can create new samples with  $D$  calls to the neural net.
- ▶ At each step we randomly choose  $u_i$  from  $N(0, 1)$  and scale and shift it to get  $\tilde{x}_i = \exp(\alpha) * u_i + \mu$ .
- ▶ At each step the  $\alpha$  and  $\mu$  come from a call to the net with  $\tilde{x}_1, \dots, \tilde{x}_{i-1}$  as inputs.

# Sampling from the inferred probability density 1

- ▶ Thus we get a mapping  $f$  from  $u \sim N(0, I)$  to  $\tilde{x} \sim \tilde{p}$ .
- ▶ The mapping is easily invertible!
- ▶  $u_i = (\tilde{x}_i - \mu) \exp(-\alpha)$
- ▶ All the  $\mu$  and  $\alpha$  that we need are available in **one** call to the net (as here we have all the net input values that we need).
- ▶ This gives  $f^{-1}$  mapping  $\tilde{p}$  to  $N(0, I)$ .

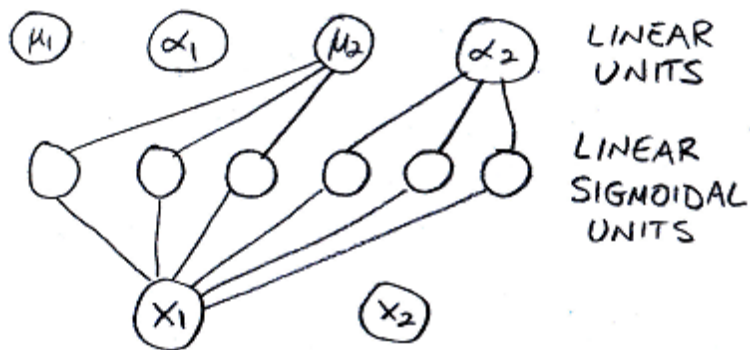
# Sampling from the inferred probability density 2

- ▶ Furthermore the autoregressive property means that the Jacobian of  $f^{-1}$  is triangular.
- ▶ Its diagonal elements are just the  $\exp(-\alpha_i)$  factors.
- ▶ So the determinant of  $Jac(f^{-1})$  is easy to calculate.


# Normalising flow

- ▶ So  $f$  maps  $N(0, I)$  to  $\tilde{p}$ .
- ▶ Such  $f$  is called a normalising flow.
- ▶ Recall how probability densities transform:  
$$\tilde{p}(x) = \mathcal{N}_{(0, I)}(f^{-1}(x)) | \det(\text{Jac}(f^{-1})) |$$
- ▶ All of these components are easy to calculate, and let us evaluate and sample from  $\tilde{p}$ .

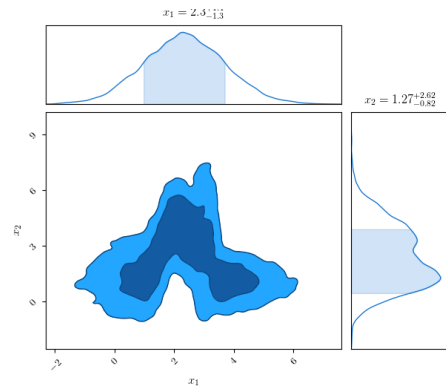
## Example - Neural net



## Example - Target distribution

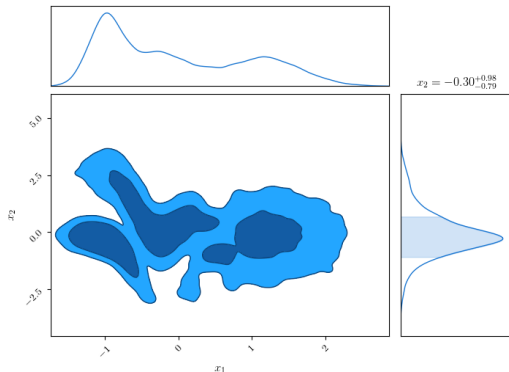


# Example - inferred distribution





Inverse image of inferred distribution - should be  $N(0, I)$



# MAF idea

- ▶ Map the original training data back through  $f^{-1}$  to get  $\{u_i\}$  that should be  $N(0, I)$  but aren't;
- ▶ Repeat the whole process! Treat the  $u$  as training data for a new neural net that is trying to model the  $u$  as a distorted  $N(0, I)$
- ▶ Be prepared to do this five or ten iterations.

# MAF - after two iterations

